

Agent Passport Control

The permissions layer for AI agents on the website side: verify which agent is acting, record what it does, and grant access to specific actions based on each company's rules.

The problem

AI agents now use SaaS products while logged in as real customers. Claude in Chrome, Perplexity Comet, ChatGPT and Gemini in Chrome drive the customer's own browser, with their login, cookies and device, so bot detection and authentication both let them straight through. Most of this isn't malicious: it's customers getting work done faster. But today a SaaS company can't answer three questions:

Who is driving?

Is this session a person, or an agent acting for them? Which agent?

What may it do?

Nobody has decided which actions an agent can take for this person.

How much does it use?

Every agent action is logged and priced as if the person did it.

How it works: verify, record, permit

01 · VERIFY Is an agent here, and which? We detect agents using the site and name them, from a signature or from how the session is driven: the agent's passport.	02 · RECORD What it does, for whom Every agent action is logged with the agent and the person it acts for, so the company sees all agent activity.	03 · PERMIT What it's allowed to do Each company decides which agents get access to which actions. That permission is the visa, checked on every action.
--	---	---

What it does

01 · VERIFY

See every agent

One script tag tells when an agent is using the site, names the product from a catalogue of 70 agents, and marks the moment a person hands the session to an agent. It reads *how* the page is driven (pointer paths, click and key timing, how text arrives), never *what* is typed or shown. Agents that sign their requests, such as OpenAI's ChatGPT agent, are verified cryptographically.

03 · PERMIT

Visas, set in three layers

- The SaaS company sets the rules: which agents may ever get which permissions on its product.
- Its customer's admin narrows them for their own team.
- The person issues the visa to their own agent, once, in plain words: *"Claude wants to view and edit invoices for 30 days."*
- Scopes are product actions (view, export, edit, invite, pay, settings), so they work for agents that click through the UI.
- Visas expire by default and revoke in one click.

05 · INSIGHTS

Know how agents use the product

Which agents use the product, which features they rely on, which customer accounts lean on them, and where they fail.

02 · RECORD

A record of who really acted

Every agent action is recorded with the agent, its tier, its visa and the person it acted for. Exportable for SOC 2 audits, regulated customers and the EU AI Act.

Claude in Chrome · for Morgan · visa V-7F3A
exported 4,210 rows · 10:44

04 · CONTROL

A decision on every sensitive action

At exports, deletes, invites, payments and settings changes, one API call checks the agent's visa against the company's rules and returns one of seven outcomes:

Admit

Slow

Request a visa

Ask the person

Reroute to API/tool

Bill

Refuse

"Ask the person" uses their passkey, which an agent can't approve. Every decision goes into the record.

06 Tiered agent access

Every session gets a trust tier, and rules are written against tiers, not one agent at a time. The stronger the proof of who the agent is, the more it can be allowed. Deleting stays people-only by default; customers can change any default.

TIER	HOW WE KNOW	DEFAULT ACCESS
Verified	Its provider signed the request, or it identified itself	Most actions allowed, with fewer challenges. Payments and settings need the person's passkey
Recognised	Behaviour matches a known product, such as Claude in Chrome or Comet	View freely. Exports need a visa. Edits, invites and sends need the person to co-sign. Payments and settings refused
Unknown automation	Clearly automated, product unknown	View only, slowed down. Everything else refused
Human	No agent detected	Normal access

07 Seats and agent pricing

Seats are priced for a person working a few hours a day. An agent on the same login can run all night and all weekend, so a team could buy two seats and have agents do the work of ten.

We split **person-hours from agent-hours on the same login**, per account, and can pass agent usage to billing: charge for it, include an allowance in each plan, or simply allow it.

One seat, one week (illustrative)

Person

~40 h

Agent, same login

up to ~140 h

Today both show up as one very busy user.

Example: who issues the visa

COMPANY B · OUR CUSTOMER

Ledgerline

Sets the rules for its accounting product: *financial records can only be edited by approved agents, with the account holder's visa. Agents never delete records or change payment details.*

COMPANY A

Acme's admin

Narrows them for Acme's workspace: *our team may use Claude in Chrome, to view, export and edit invoices.*

ACME EMPLOYEE

Morgan

Issues the visa: approves "*Claude wants to view and edit invoices for 30 days.*" Claude is now tied to Morgan's account, and every action it takes is checked against that visa.

MORGAN'S AGENT

Claude in Chrome

Does the work. When a malicious email tells it to export the customer list, that's outside its visa: **refused**, and logged against Claude, acting for Morgan. An agent outside these rules, such as Comet here, never gets a visa at all.

Who it's for

B2B SaaS with sensitive data (CRM, finance, HR, analytics), seat-priced products and regulated SaaS. **Security** gets control and an audit trail, **product** gets visibility and a smooth path for good agents, and **pricing / RevOps** gets agent usage and billing.

Where we are

BUILT

Sensor with a 70-agent catalogue and hand-off detection. Console with live sessions, agents, person vs agent hours and the entry log.

IN PROGRESS

A benchmark of accuracy against real agents and real people, including screen-reader and dictation users.

NEXT

Visas and the consent prompt, the decision API, enforcement, and billing integration.